

CSCI544 Assignment 1

Vorapoom Thirapatarapong

1. Dataset Preparation

The data is firstly downloaded and decompressed. Then, the Pandas' read_csv function is used to read the data into dataframe. Note that there are a few lines with incorrect formats, and these lines are skipped. After that, only the review and star rating columns are kept, the ratings are grouped into 3 classes as specified, and 20,000 data points from each class are randomly selected for training in the future.

2. Data Cleaning

Few steps are done during the data cleaning process. Firstly, the reviews are converted to lowercase, and the HTML and URLs from the reviews are removed. Then, the contractions are performed, following by removing the non-alphabetical characters (contraction step is done prior to removing non-alphabetical characters as some punctuations are needed to perform contractions correctly). Finally, extra spaces are removed. The average length of reviews after data cleaning reduces slightly (<5%)

3. Data Preprocessing

NLTK package is used in this step to remove stop words from the reviews and perform lemmatization. More detail on the lemmatization step is that the POS tagging is also used when lemmatizing a word. If the POS is either a verb or an adjective, its POS is passed to the lemmatize function. Else, the default lemmatizing parameters are used. This step is done to get a more standardized vocabulary after the process. The average length of reviews after data preprocessing reduces by around 40%.

4. Feature Extraction

The sklearn's TfidfVectorizer function is used to extract features from the reviews from the previous step. The parameter max_features is used and set to 10,000 features. This helps reduce training time and doesn't affect the performance.

5. Perceptron

The sklearn's Perceptron library is used. With the parameter penalty set to elasticnet, it yields the average precision score across 3 classes at 66.55%

6. SVM

The sklearn's LinearSVC library is used. With the parameter multiclass set to ovr and max_iter set to 50,000, it yields the average precision score across 3 classes at 65.42%

7. Logistic Regression

The sklearn's LogisticRegression library is used. With the parameter multiclass set to ovr and max_iter set to 5,000, it yields the average precision score across 3 classes at 67.29%

8. Multinomial Naïve Bayes

The sklearn's MultinomialNB library is used. With the default parameter, it yields the average precision score across 3 classes at 65.81%

9. Others

- References:
 - i. <https://stackoverflow.com/questions/45999415/removing-html-tags-in-pandas>
 - ii. <https://stackoverflow.com/questions/11331982/how-to-remove-any-url-within-a-string-in-python>
 - iii. <https://www.machinelearningplus.com/nlp/lemmatization-examples-python/>
- Libraries used
 - i. Python version = 3.8.8
 - ii. beautifulsoup4==4.11.1
 - iii. contractions==0.1.73
 - iv. nltk==3.6.1
 - v. numpy==1.24.1
 - vi. pandas==1.5.3
 - vii. scikit_learn==1.2.1

```
In [1]: import pandas as pd
import numpy as np
import nltk
nltk.download('wordnet')
nltk.download('omw-1.4')
nltk.download('averaged_perceptron_tagger')
import re
from bs4 import BeautifulSoup
import contractions
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from sklearn.linear_model import Perceptron
from sklearn.svm import LinearSVC
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
```

```
/opt/anaconda3/lib/python3.8/site-packages/scipy/__init__.py:138: UserWarning:
A NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy (de
tected version 1.24.1)
```

```
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion} is re
quired for this version of ")
```

```
[nltk_data] Downloading package wordnet to /Users/boom/nltk_data...
```

```
[nltk_data] Package wordnet is already up-to-date!
```

```
[nltk_data] Downloading package omw-1.4 to /Users/boom/nltk_data...
```

```
[nltk_data] Package omw-1.4 is already up-to-date!
```

```
[nltk_data] Downloading package averaged_perceptron_tagger to
```

```
[nltk_data] /Users/boom/nltk_data...
```

```
[nltk_data] Package averaged_perceptron_tagger is already up-to-
```

```
[nltk_data] date!
```

```
In [2]: # ! pip install bs4 # in case you don't have it installed
# ! pip install contractions
```

```
# Dataset: https://s3.amazonaws.com/amazon-reviews-pds/tsv/amazon_reviews_us_
```

1. Dataset Preparation

Read Data

```
In [3]: df = pd.read_csv('data.tsv', sep='\t', on_bad_lines='skip')
print(df.head())
```

```
<ipython-input-3-a057dc4512f1>:1: DtypeWarning: Columns (7) have mixed types.
Specify dtype option on import or set low_memory=False.
```

```
df = pd.read_csv('data.tsv', sep='\t', on_bad_lines='skip')
```

	marketplace	customer_id	review_id	product_id	product_parent	\
0	US	1797882	R3I2DHQBR577SS	B001ANOOOE	2102612	
1	US	18381298	R1QNE9NQFJC2Y4	B0016J22EQ	106393691	
2	US	19242472	R3LIDG2Q4LJBAO	B00HU6UQAG	375449471	
3	US	19551372	R3KSZHPAEVPEAL	B002HWS7RM	255651889	
4	US	14802407	RAI2OIG50KZ43	B00SM99KWU	116158747	

	product_title	product_category	\
0	The Naked Bee Vitmin C Moisturizing Sunscreen ...	Beauty	
1	Alba Botanica Sunless Tanning Lotion, 4 Ounce	Beauty	
2	Elysee Infusion Skin Therapy Elixir, 2oz.	Beauty	
3	Diane D722 Color, Perm And Conditioner Process...	Beauty	

4 Biore UV Aqua Rich Watery Essence SPF50+/PA+++... Beauty

	star_rating	helpful_votes	total_votes	vine	verified_purchase	\
0	5	0.0	0.0	N		Y
1	5	0.0	0.0	N		Y
2	5	0.0	0.0	N		Y
3	5	0.0	0.0	N		Y
4	5	0.0	0.0	N		Y

	review_headline	\
0	Five Stars	
1	Thank you Alba Bontanica!	
2	Five Stars	
3	GOOD DEAL!	
4	this soaks in quick and provides a nice base f...	

	review_body	review_date
0	Love this, excellent sun block!!	2015-08-31
1	The great thing about this cream is that it do...	2015-08-31
2	Great Product, I'm 65 years old and this is al...	2015-08-31
3	I use them as shower caps & conditioning caps....	2015-08-31
4	This is my go-to daily sunblock. It leaves no ...	2015-08-31

Keep Reviews and Ratings

```
In [4]: df = df.loc[:, ['review_body', 'star_rating']]
df['star_rating'] = pd.to_numeric(df['star_rating'], errors='coerce')
df = df[-df['star_rating'].isna()]
df['star_rating'] = df['star_rating'].astype(int)
```

```
In [5]: df = df.loc[:, ['review_body', 'star_rating']]
print(df.head())
```

	review_body	star_rating
0	Love this, excellent sun block!!	5
1	The great thing about this cream is that it do...	5
2	Great Product, I'm 65 years old and this is al...	5
3	I use them as shower caps & conditioning caps....	5
4	This is my go-to daily sunblock. It leaves no ...	5

Group ratings to 3 classes

```
In [6]: mapping = {1: 1, 2: 1, 3: 2, 4: 3, 5: 3}
df = df.replace({'star_rating': mapping})
print(df.head())
```

	review_body	star_rating
0	Love this, excellent sun block!!	3
1	The great thing about this cream is that it do...	3
2	Great Product, I'm 65 years old and this is al...	3
3	I use them as shower caps & conditioning caps....	3
4	This is my go-to daily sunblock. It leaves no ...	3

We form three classes and select 20000 reviews randomly from each class.

```
In [7]: # drop duplicates
df = df.drop_duplicates()
```

```
In [8]: r_state = 555
list_sample = [df[df.star_rating == 1].sample(n=20000, random_state=r_state),
               df[df.star_rating == 2].sample(n=20000, random_state=r_state),
               df[df.star_rating == 3].sample(n=20000, random_state=r_state)]
df_sample = pd.concat(list_sample)
print(df_sample.head())
```

	review_body	star_rating
1294234	Handle provided had some rusted internal threa...	1
4186939	MY WIFE SAYS THAT THIS REALLY WASNT USEFUL FOR...	1
4493244	This came with no glue, and the gel was hard a...	1
3048924	I bought this product because of all the 5-sta...	1
4149821	i got this product expecting the wow factor, b...	1

```
In [9]: print(df_sample.groupby('star_rating').count())
```

star_rating	review_body
1	20000
2	20000
3	20000

2. Data Cleaning

Convert to lowercase

```
In [10]: df_clean = df_sample.copy()
df_clean.columns = ['review', 'stars']

df_clean['review'] = df_clean['review'].apply(str.lower)
print(df_clean.head())
```

	review	stars
1294234	handle provided had some rusted internal threa...	1
4186939	my wife says that this really wasnt useful for...	1
4493244	this came with no glue, and the gel was hard a...	1
3048924	i bought this product because of all the 5-sta...	1
4149821	i got this product expecting the wow factor, b...	1

remove the HTML and URLs from the reviews

```
In [11]: # remove HTML
df_clean['review'] = df_clean['review'].str.replace(r'<[^>]*>', '', regex=True)

# remove URLs
def remove_urls(text):
    return re.sub(r'((?i)\b((?:https?://|www\d{0,3}[.]|[a-z0-9.\-]+[.][a-z]{2,4})\b|[a-z0-9.\-]+[.][a-z]{2,4})', '', text)

df_clean['review'] = df_clean['review'].apply(remove_urls)
print(df_clean.head())
```

	review	stars
1294234	handle provided had some rusted internal threa...	1
4186939	my wife says that this really wasnt useful for...	1
4493244	this came with no glue, and the gel was hard a...	1
3048924	i bought this product because of all the 5-sta...	1
4149821	i got this product expecting the wow factor, b...	1

perform contractions on the reviews

```
In [12]: def perform_contractions(text):
          return ' '.join([contractions.fix(word) for word in text.split()])

df_clean['review'] = df_clean['review'].apply(perform_contractions)
print(df_clean.head())
```

	review	stars
1294234	handle provided had some rusted internal threa...	1
4186939	my wife says that this really was not useful f...	1
4493244	this came with no glue, and the gel was hard a...	1
3048924	i bought this product because of all the 5-sta...	1
4149821	i got this product expecting the wow factor, b...	1

remove non-alphabetical characters

```
In [13]: def remove_non_alpha_chars(text):
          return re.sub(r'^a-zA-Z0-9\s', ' ', text)

df_clean['review'] = df_clean['review'].apply(remove_non_alpha_chars)
print(df_clean.head())
```

	review	stars
1294234	handle provided had some rusted internal threa...	1
4186939	my wife says that this really was not useful f...	1
4493244	this came with no glue and the gel was hard a...	1
3048924	i bought this product because of all the 5 sta...	1
4149821	i got this product expecting the wow factor b...	1

remove extra spaces

```
In [14]: def remove_extra_spaces(text):
          return re.sub(r'\s+', ' ', text)

df_clean['review'] = df_clean['review'].apply(remove_extra_spaces)
print(df_clean.head())
```

	review	stars
1294234	handle provided had some rusted internal threa...	1
4186939	my wife says that this really was not useful f...	1
4493244	this came with no glue and the gel was hard as...	1
3048924	i bought this product because of all the 5 sta...	1
4149821	i got this product expecting the wow factor bu...	1

```
In [15]: # printing average lengths before/after data cleaning
print('Average length of reviews before and after data cleaning: ', \
      df_sample['review_body'].str.len().mean(), ', ', df_clean['review'].str
```

Average length of reviews before and after data cleaning: 281.41571666666664, 273.19905

3. Pre-processing

Remove the stop words

```
In [16]: stop_words = set(stopwords.words('english'))

def remove_stop_words(text):
    return ' '.join([word for word in text.split() if word not in stop_words])
```

```
df_preproc = df_clean.copy()
df_preproc['review'] = df_preproc['review'].apply(remove_stop_words)
print(df_preproc.head())
```

	review	stars
1294234	handle provided rusted internal threading incl...	1
4186939	wife says really useful nail polish plate rubb...	1
4493244	came glue gel hard rock never buying product d...	1
3048924	bought product 5 star ratings received wonderi...	1
4149821	got product expecting wow factor found priced ...	1

Perform lemmatization

In [17]:

```
lemmatizer = WordNetLemmatizer()

def perform_lemmatization(lemmatizer, text):
    lemmatized_list = []
    for word, pos_tag in nltk.pos_tag(text.split()):
        if pos_tag.startswith('V'):
            lemmatized_list.append(lemmatizer.lemmatize(word, 'v'))
        elif pos_tag.startswith('J'):
            lemmatized_list.append(lemmatizer.lemmatize(word, 'a'))
        else:
            lemmatized_list.append(lemmatizer.lemmatize(word))
    return ' '.join(lemmatized_list)

df_lemma = df_preproc.copy()
df_lemma['review'] = df_lemma['review'].apply(lambda x: perform_lemmatization(x, lemmatizer))
print(df_lemma.head())
```

	review	stars
1294234	handle provide rust internal threading include...	1
4186939	wife say really useful nail polish plate rub s...	1
4493244	come glue gel hard rock never buy product diss...	1
3048924	buy product 5 star rating receive wonder produ...	1
4149821	get product expect wow factor find price size ...	1

In [18]:

```
# printing average lengths before/after data preprocessing
print('Average length of reviews before and after data preprocessing: ', \
      df_clean['review'].str.len().mean(), ', ', df_lemma['review'].str.len().mean())
```

Average length of reviews before and after data preprocessing: 273.19905, 158.77876666666666

4. TF-IDF Feature Extraction

In [19]:

```
vectorizer = TfidfVectorizer(max_features=10000)

df_model = df_lemma.copy()
tfidf = vectorizer.fit_transform(df_model['review'])
df_X = pd.DataFrame(tfidf.toarray(), columns=vectorizer.get_feature_names())
```

In [20]:

```
print(df_X.head())
df_y = df_model['stars'].reset_index(drop=True)
df_y = df_y.astype(int)
print(df_y.head())
```

	00	000	01	02	03	04	05	06	07	08	...	ziploc	ziplock	\
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	

2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0

	zipper	zirconium	zit	zits	zombie	zone	zoom	zoya
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

[5 rows x 10000 columns]

```
0    1
1    1
2    1
3    1
4    1
```

Name: stars, dtype: int64

Split training/testing set - 80/20

In [21]:

```
X_train, X_test, y_train, y_test = train_test_split(df_X, df_y, test_size=0.2)
print(X_train, X_test, y_train, y_test)
```

\	00	000	01	02	03	04	05	06	07	08	...	ziploc	ziplock
5521	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
38538	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
30253	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
45733	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
2593	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
...
39184	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
18365	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
33132	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
22889	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
10375	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0

	zipper	zirconium	zit	zits	zombie	zone	zoom	zoya
5521	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
38538	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
30253	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
45733	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2593	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
...
39184	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
18365	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
33132	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
22889	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
10375	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

[48000 rows x 10000 columns]	00	000	01	02	03	04	05	06	0
7 08 ... ziploc ziplock \									
31788	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
16832	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
8267	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
55958	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
11472	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
...
5042	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
35683	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
52589	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
6437	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
33910	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

	zipper	zirconium	zit	zits	zombie	zone	zoom	zoya
31788	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
16832	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

8267	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
55958	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
11472	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
...
5042	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
35683	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
52589	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
6437	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
33910	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

```
[12000 rows x 10000 columns] 5521      1
38538      2
30253      2
45733      3
2593       1
      ..
39184      2
18365      1
33132      2
22889      2
10375      1
Name: stars, Length: 48000, dtype: int64 31788      2
16832      1
8267       1
55958      3
11472      1
      ..
5042       1
35683      2
52589      3
6437       1
33910      2
Name: stars, Length: 12000, dtype: int64
```

5. Perceptron

```
In [22]: def report_scores(clf, X_test, y_test, clf_name):
          y_pred = clf.predict(X_test)
          dict_report = classification_report(y_test, y_pred, output_dict=True)
          key_list = ['1', '2', '3', 'weighted avg']
          report_name_list = ['(Class 1)', '(Class 2)', '(Class 3)', '(Average)']
          for i, k in enumerate(key_list):
              temp_report = dict_report[k]
              precision = str(temp_report['precision'])
              recall = str(temp_report['recall'])
              f1_score = str(temp_report['f1-score'])
              print('Precision, Recall, and f1-score for the testing split for ' +
                    precision + ',' + recall + ',' + f1_score)
```

```
In [23]: clf_perceptron = Perceptron(random_state=r_state, penalty='elasticnet')
          clf_perceptron.fit(X_train, y_train)
```

```
Out[23]: Perceptron(penalty='elasticnet', random_state=555)
```

```
In [24]: report_scores(clf_perceptron, X_test, y_test, 'Perceptron')
```

```
Precision, Recall, and f1-score for the testing split for Perceptron (Class
1):  0.6639344262295082,0.6075,0.6344647519582245
Precision, Recall, and f1-score for the testing split for Perceptron (Class
2):  0.4931712191872085,0.74025,0.5919632147141144
```

Precision, Recall, and f1-score for the testing split for Perceptron (Class 3): 0.8394691780821918,0.49025,0.6190025252525253
 Precision, Recall, and f1-score for the testing split for Perceptron (Average): 0.6655249411663028,0.6126666666666667,0.6151434973082881

6. SVM

```
In [25]: clf_SVC = LinearSVC(random_state=r_state, multi_class='ovr', dual=True, max_iter=50000)
         clf_SVC.fit(X_train, y_train)
```

```
Out[25]: LinearSVC(max_iter=50000, random_state=555)
```

```
In [26]: report_scores(clf_SVC, X_test, y_test, 'SVM')
```

Precision, Recall, and f1-score for the testing split for SVM (Class 1): 0.6718175128771156,0.68475,0.6782221121703603
 Precision, Recall, and f1-score for the testing split for SVM (Class 2): 0.5754189944134078,0.54075,0.5575460755251966
 Precision, Recall, and f1-score for the testing split for SVM (Class 3): 0.7154178674351584,0.74475,0.7297893189612935
 Precision, Recall, and f1-score for the testing split for SVM (Average): 0.6542181249085607,0.65675,0.6551858355522835

7. Logistic Regression

```
In [27]: clf_logreg = LogisticRegression(random_state=r_state, multi_class='ovr', max_iter=5000)
         clf_logreg.fit(X_train, y_train)
```

```
Out[27]: LogisticRegression(max_iter=5000, multi_class='ovr', random_state=555)
```

```
In [28]: report_scores(clf_logreg, X_test, y_test, 'Logistic Regression')
```

Precision, Recall, and f1-score for the testing split for Logistic Regression (Class 1): 0.6900655817342726,0.71025,0.7000123198225946
 Precision, Recall, and f1-score for the testing split for Logistic Regression (Class 2): 0.5955862802446158,0.56,0.5772452003607782
 Precision, Recall, and f1-score for the testing split for Logistic Regression (Class 3): 0.7331392527899078,0.7555,0.744151686776656
 Precision, Recall, and f1-score for the testing split for Logistic Regression (Average): 0.6729303715895987,0.67525,0.6738030689866763

8. Naive Bayes

```
In [29]: clf_nb = MultinomialNB()
         clf_nb.fit(X_train, y_train)
```

```
Out[29]: MultinomialNB()
```

```
In [30]: report_scores(clf_nb, X_test, y_test, 'Naive Bayes')
```

Precision, Recall, and f1-score for the testing split for Naive Bayes (Class 1): 0.6969292389853138,0.6525,0.6739832149774048
 Precision, Recall, and f1-score for the testing split for Naive Bayes (Class 2): 0.5639562529719448,0.593,0.5781135754326103
 Precision, Recall, and f1-score for the testing split for Naive Bayes (Class 3): 0.7331392527899078,0.7555,0.744151686776656

3): 0.7135095085206223,0.72225,0.7178531494595601

Precision, Recall, and f1-score for the testing split for Naive Bayes (Average): 0.6581316668259604,0.6559166666666667,0.6566499799565251

References

- <https://stackoverflow.com/questions/45999415/removing-html-tags-in-pandas>
- <https://stackoverflow.com/questions/11331982/how-to-remove-any-url-within-a-string-in-python>
- <https://www.machinelearningplus.com/nlp/lemmatization-examples-python/>