

# CSCI-544 HOMEWORK 3

Name: Vorapoom Thirapatarapong

In [1]:

```
# packages
import pandas as pd
import numpy as np
import time
import copy
import nltk
nltk.download('wordnet')
nltk.download('omw-1.4')
nltk.download('averaged_perceptron_tagger')
import re
from bs4 import BeautifulSoup
import contractions
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import Perceptron
from sklearn.svm import LinearSVC
import gensim.downloader as api
import gensim.models
from gensim.test.utils import datapath
from gensim import utils
from numpy import dot
from numpy.linalg import norm
import torch
import torch.utils.data as data_utils
import torch.nn as nn
import torch.nn.functional as F
```

```
/opt/anaconda3/lib/python3.8/site-packages/scipy/__init__.py:138: UserWarning:
A NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy (de
tected version 1.24.1)
```

```
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion} is re
quired for this version of ")
```

```
[nltk_data] Downloading package wordnet to /Users/boom/nltk_data...
```

```
[nltk_data] Package wordnet is already up-to-date!
```

```
[nltk_data] Downloading package omw-1.4 to /Users/boom/nltk_data...
```

```
[nltk_data] Package omw-1.4 is already up-to-date!
```

```
[nltk_data] Downloading package averaged_perceptron_tagger to
```

```
[nltk_data] /Users/boom/nltk_data...
```

```
[nltk_data] Package averaged_perceptron_tagger is already up-to-
```

```
[nltk_data] date!
```

In [2]:

```
# parameters
r_state = 555
sample_size = 20000
test_ratio = 0.2
torch.manual_seed(r_state)
```

Out[2]: <torch.\_C.Generator at 0x7f820f478cf0>

# 1. Dataset Generation

## Load data + Initial cleaning

```
In [3]: # read dataset
df = pd.read_csv('data.tsv', sep='\t', on_bad_lines='skip')

<ipython-input-3-f74a6e2e7113>:2: DtypeWarning: Columns (7) have mixed types.
Specify dtype option on import or set low_memory=False.
    df = pd.read_csv('data.tsv', sep='\t', on_bad_lines='skip')

In [4]: # clean invalid rows & format data types
df = df.loc[:, ['review_body', 'star_rating']]
df['star_rating'] = pd.to_numeric(df['star_rating'], errors='coerce')
df = df[~df['star_rating'].isna()]
df = df[~df['review_body'].isna()]
df['star_rating'] = df['star_rating'].astype(int)

In [5]: # columns selection
df = df.loc[:, ['review_body', 'star_rating']]

In [6]: # drop duplicates
df = df.drop_duplicates()

In [7]: # group ratings to 3 classes
mapping = {1: 0, 2: 0, 3: 1, 4: 2, 5: 2}
df = df.replace({'star_rating': mapping})

In [8]: # save a copy of the ,000 sample reviews for training custom Word2Vec
list_w2v = [df[df.star_rating == 0].sample(n=20000, random_state=r_state),
            df[df.star_rating == 1].sample(n=20000, random_state=r_state),
            df[df.star_rating == 2].sample(n=20000, random_state=r_state)]
df_w2v = pd.concat(list_w2v)['review_body']

In [9]: # sample a balanced dataset of 60k reviews
list_sample = [df[df.star_rating == 0].sample(n=sample_size, random_state=r_s
            df[df.star_rating == 1].sample(n=sample_size, random_state=r_stat
            df[df.star_rating == 2].sample(n=sample_size, random_state=r_stat
df_sample = pd.concat(list_sample)

In [10]: # columns renaming
df_sample.columns = ['review', 'stars']

In [11]: # verify number of samples
print(df_sample.groupby('stars').count())
```

	review
stars	
0	20000
1	20000
2	20000

## Text cleaning

```
In [12]:
def remove_urls(text):
    return re.sub(r'((?i)\b((?:https?://|www\d{0,3}[.]|[a-z0-9.\-]+[.])[a-z]{2,5})', '', text)

def perform_contractions(text):
    return ' '.join([contractions.fix(word) for word in text.split()])

def remove_non_alpha_chars(text):
    return re.sub(r'^a-zA-Z0-9\s', ' ', text)

def remove_extra_spaces(text):
    return re.sub(r'\s+', ' ', text)
```

```
In [13]:
df_clean = df_sample.copy()

# lowercase
df_clean['review'] = df_clean['review'].apply(str.lower)

# remove HTML
df_clean['review'] = df_clean['review'].str.replace(r'<[^>]*>', '', regex=True)

# remove URLs
df_clean['review'] = df_clean['review'].apply(remove_urls)

# contractions
df_clean['review'] = df_clean['review'].apply(performance_contractions)

# non-alphabet chars
df_clean['review'] = df_clean['review'].apply(remove_non_alpha_chars)

# extra spaces
df_clean['review'] = df_clean['review'].apply(remove_extra_spaces)
```

## Text Preprocessing

```
In [14]:
def remove_stop_words(text):
    return ' '.join([word for word in text.split() if word not in stop_words])

def perform_lemmatization(lemmatizer, text):
    lemmatized_list = []
    for word, pos_tag in nltk.pos_tag(text.split()):
        if pos_tag.startswith('V'):
            lemmatized_list.append(lemmatizer.lemmatize(word, 'v'))
        elif pos_tag.startswith('J'):
            lemmatized_list.append(lemmatizer.lemmatize(word, 'a'))
        else:
            lemmatized_list.append(lemmatizer.lemmatize(word))
    return ' '.join(lemmatized_list)
```

```
In [15]:
# remove stop words
stop_words = set(stopwords.words('english'))

df_preproc = df_clean.copy()
df_preproc['review'] = df_preproc['review'].apply(remove_stop_words)

# lemmatization
lemmatizer = WordNetLemmatizer()
```

```
df_lemma = df_preproc.copy()
df_lemma['review'] = df_lemma['review'].apply(lambda x: perform_lemmatization
```

## Split Training / Testing

```
In [16]: # define X and y
df_X = df_lemma['review'].reset_index(drop=True)
df_y = df_lemma['stars'].reset_index(drop=True)

# split training/testing = 80/20
X_train, X_test, y_train, y_test = train_test_split(df_X, df_y, test_size=tes
```

## TF-IDF features extraction

```
In [17]: vectorizer = TfidfVectorizer(max_features=10000)

df_X_tfidf_pre = pd.concat([X_train, X_test])
tfidf = vectorizer.fit_transform(df_X_tfidf_pre)
df_X_tfidf = pd.DataFrame(tfidf.toarray(), columns=vectorizer.get_feature_names())

X_train_tfidf = df_X_tfidf[:int(3*sample_size*(1-test_ratio))]
X_test_tfidf = df_X_tfidf[int(3*sample_size*(1-test_ratio)):]
```

# 2. Word Embedding

## 2(a). Pre-trained Word2Vec

```
In [18]: wv = api.load('word2vec-google-news-300')
```

```
In [19]: # Semantics similarity example #1

a = wv['boat'] - wv['water'] + wv['air']
b = wv['plane']

cos_sim = dot(a, b)/(norm(a)*norm(b))
print(cos_sim)
```

0.5111032

```
In [20]: # Semantics similarity example #2

a = wv['sea']
b = wv['ocean']

cos_sim = dot(a, b)/(norm(a)*norm(b))
print(cos_sim)
```

0.76435417

```
In [21]: # Semantics similarity example #3

a = wv['vegetable']
b = wv['star']
```

```
cos_sim = dot(a, b)/(norm(a)*norm(b))
print(cos_sim)
```

0.023547666

## 2(b). Custom Word2Vec

```
In [22]: # create a corpus using 500,000 reviews
class MyCorpus:
    """An iterator that yields sentences (lists of str)."""

    def __iter__(self):
        for index, row in df_w2v.iteritems():
            # print(row)
            yield utils.simple_preprocess(row)
```

```
In [23]: sentences = MyCorpus()
model = gensim.models.Word2Vec(sentences, min_count=9, vector_size=300, window=5)
```

<ipython-input-22-2c0d12710906>:6: FutureWarning: iteritems is deprecated and will be removed in a future version. Use .items instead.

```
for index, row in df_w2v.iteritems():
```

```
In [24]: # Semantics similarity example #1

a = model.wv['boat'] - model.wv['water'] + model.wv['air']
b = model.wv['plane']

cos_sim = dot(a, b)/(norm(a)*norm(b))
print(cos_sim)
```

-0.046374116

```
In [25]: # Semantics similarity example #2

a = model.wv['sea']
b = model.wv['ocean']

cos_sim = dot(a, b)/(norm(a)*norm(b))
print(cos_sim)
```

0.59367234

```
In [26]: # Semantics similarity example #3

a = model.wv['vegetable']
b = model.wv['star']

cos_sim = dot(a, b)/(norm(a)*norm(b))
print(cos_sim)
```

0.09536163

What do you conclude from comparing vectors generated by yourself and the pretrained model?

The dimensions of the vectors from both pretrained(word2vec-google-news-300) and custom Word2Vec models are the same which is 300. However, the pretrained model is able to encode the semantics similarity much better -

when looking at two similar words, the pretrained model is able to give higher similarity score & when looking at 2 words with no relationship at all, the pretrained model is able to give a slightly lower similarity score. This might be because that the training samples of the pretrained model is much larger which makes it generalize better.

Which of the Word2Vec models seems to encode semantic similarities between words better?

From the 3 examples chosen, it seems like the pretrained Word2Vec model is able to encode the semantic similarities better. For the first 2 examples, the two objects are highly related which should give higher similarity value, and the pretrained model is giving higher similarities. For the third example which is the two objects that are not related, both models seem to satisfy the condition where they give relatively low similarity scores between the two objects. The reason for better semantic similarity encoding for the pretrained model might be a much larger size of training samples.

### 3. Simple Models: Single perceptron & SVM

In [27]:

```
# Functions for W2V Feature extraction
def extract_features_w2v(data, max_vec_concat, max_vec_rnn):
    data = list(data)
    list_output_avg = list()
    list_output_concat = list()
    df_output_concat = pd.DataFrame(columns=['c'+str(i) for i in range(1, 300)])
    np_output_3d = np.zeros((len(data), max_vec_rnn, 300), dtype=np.float32)
    x = 0
    time_bf, time_start = time.time(), time.time()

    for i, instance in enumerate(data):

        cnt = 0
        list_temp = []
        list_temp2 = []
        for word in instance.split():
            try:
                vec = wv[word]
                list_temp2.append(list(vec))
                list_temp = list_temp + list(vec)
                if cnt < 20:
                    np_output_3d[i, cnt, :] = vec
                    cnt += 1
            except:
                pass

        temp_avg = np.mean(list_temp2, axis=0)
        list_output_avg.append(temp_avg)
        list_temp = (list_temp + max_vec_concat*300*[0.0])[:max_vec_concat*300]
        list_output_concat.append(list_temp)

        x += 1
        if x % 2000 == 0:
            print(x, time.time() - time_bf)
            time_bf = time.time()

    list_output_avg = [x if isinstance(x, np.ndarray) else np.zeros(300) for
```

```
list_output_concat = [x if len(x) > 0 else [0.0]*3000 for x in list_output]

df_output_avg = pd.DataFrame(list_output_avg, columns=['c'+str(i) for i in range(1, 3000)])
df_output_concat = pd.DataFrame(list_output_concat, columns=['c'+str(i) for i in range(1, 3000)])

print(time.time() - time_start)

return df_output_avg, df_output_concat, np_output_3d
```

### 3(a). Single perceptron

In [28]:

```
# extract Word2Vec features
X_train_w2v, X_train_w2v_concat, X_train_3d = extract_features_w2v(X_train, 10, 2)
X_test_w2v, X_test_w2v_concat, X_test_3d = extract_features_w2v(X_test, 10, 2)

2000 5.053712844848633
/opt/anaconda3/lib/python3.8/site-packages/numpy/core/fromnumeric.py:3464: RuntimeWarning: Mean of empty slice.
  return _methods._mean(a, axis=axis, dtype=dtype,
/opt/anaconda3/lib/python3.8/site-packages/numpy/core/_methods.py:192: RuntimeWarning: invalid value encountered in scalar divide
  ret = ret.dtype.type(ret / rcount)
4000 5.233086109161377
6000 5.562352895736694
8000 6.524694919586182
10000 6.205254793167114
12000 5.549623250961304
14000 6.020545959472656
16000 5.951589822769165
18000 4.73419976234436
20000 5.9541168212890625
22000 4.642481088638306
24000 5.937429666519165
26000 5.508419036865234
28000 6.115911245346069
30000 5.167878150939941
32000 5.232169151306152
34000 5.269054889678955
36000 5.526584148406982
38000 6.3384950160980225
40000 5.78859806060791
42000 6.483400821685791
44000 4.395051002502441
46000 5.530379056930542
48000 5.462535858154297
192.25431990623474
/opt/anaconda3/lib/python3.8/site-packages/numpy/core/fromnumeric.py:3464: RuntimeWarning: Mean of empty slice.
  return _methods._mean(a, axis=axis, dtype=dtype,
/opt/anaconda3/lib/python3.8/site-packages/numpy/core/_methods.py:192: RuntimeWarning: invalid value encountered in scalar divide
  ret = ret.dtype.type(ret / rcount)
2000 5.040863037109375
4000 4.650351285934448
6000 6.271064043045044
8000 6.033645153045654
10000 5.058319807052612
12000 4.746193885803223
43.062626123428345
```

In [29]:

```
X_train_w2v.fillna(0, inplace=True)
X_train_w2v_concat.fillna(0, inplace=True)
X_test_w2v.fillna(0, inplace=True)
X_test_w2v_concat.fillna(0, inplace=True)
```

```
In [30]: # TFIDF features
clf_perceptron = Perceptron(random_state=r_state, penalty='elasticnet')
clf_perceptron.fit(X_train_tfidf, y_train)
acc_tfidf_perceptron = clf_perceptron.score(X_test_tfidf, y_test)
print('Accuracy (TFIDF, Perceptron) =', acc_tfidf_perceptron)
```

Accuracy (TFIDF, Perceptron) = 0.6088333333333333

```
In [31]: # Word2Vec features
clf_perceptron = Perceptron(random_state=r_state, penalty='elasticnet')
clf_perceptron.fit(X_train_w2v, y_train)
acc_w2v_perceptron = clf_perceptron.score(X_test_w2v, y_test)
print('Accuracy (W2V, Perceptron) =', acc_w2v_perceptron)
```

Accuracy (W2V, Perceptron) = 0.5390833333333334

### 3(b). SVM

```
In [32]: # TFIDF features
clf_SVC = LinearSVC(random_state=r_state, multi_class='ovr', dual=True, max_i
clf_SVC.fit(X_train_tfidf, y_train)
acc_tfidf_svm = clf_SVC.score(X_test_tfidf, y_test)
print('Accuracy (TFIDF, SVM) =', acc_tfidf_svm)
```

Accuracy (TFIDF, SVM) = 0.65425

```
In [33]: # Word2Vec features
clf_SVC = LinearSVC(random_state=r_state, multi_class='ovr', dual=True, max_i
clf_SVC.fit(X_train_w2v, y_train)
acc_w2v_svm = clf_SVC.score(X_test_w2v, y_test)
print('Accuracy (W2V, SVM) =', acc_w2v_svm)
```

Accuracy (W2V, SVM) = 0.62775

What do you conclude from comparing performances for the models trained using the two different feature types (TF-IDF and your trained Word2Vec features)?

The accuracy values from the models that use Word2Vec as input features are much lower using perceptron and slightly lower using SVM, meaning that word-level feature extraction in this case is less meaningful for model training than using a document-level features. My assumption on this behavior is that using average Word2Vec as features are not telling anything about the sentence as a whole - the effect of the keyword that might tell the sentiment got diluted from averaging all the words / it does not consider the sequence of the words or the long-term dependencies / etc.

## 4. Feedforward Neural Networks

```
In [34]: # Converting data to tensors
X_train_tensor = torch.from_numpy(X_train_w2v.to_numpy().astype(np.float32))
X_test_tensor = torch.from_numpy(X_test_w2v.to_numpy().astype(np.float32))

y_train_tensor = torch.tensor(y_train.values)
y_test_tensor = torch.tensor(y_test.values)
```



```
# Passing to DataLoader
train_tensor = data_utils.TensorDataset(X_train_tensor, y_train_tensor)
train_loader = data_utils.DataLoader(train_tensor, batch_size=10, shuffle=True)

test_tensor = data_utils.TensorDataset(X_test_tensor, y_test_tensor)
test_loader = data_utils.DataLoader(test_tensor, batch_size=10, shuffle=True)
```

In [35]:

```
class FNN(nn.Module):
    def __init__(self, n_features):
        super(FNN, self).__init__()
        # number of hidden nodes in each layer
        hidden_1 = 100
        hidden_2 = 10

        self.fc1 = nn.Linear(n_features, hidden_1)
        self.fc2 = nn.Linear(hidden_1, hidden_2)
        self.fc3 = nn.Linear(hidden_2, 3)
        self.dropout = nn.Dropout(0.2)

    def forward(self, x):
        # add hidden layer, with relu activation function
        x = F.relu(self.fc1(x))
        # add hidden layer, with relu activation function
        x = F.relu(self.fc2(x))
        # add output layer
        x = self.fc3(x)
        return x
```

In [36]:

```
def get_acc(y_pred, y_true):
    max_scores, max_idx_class = y_pred.max(dim=1)
    num_match = (y_true == max_idx_class).sum().item()
    acc = num_match / y_true.size(0)
    return num_match, acc
```

In [37]:

```
def train(n_epochs, model, train_loader, valid_loader, optimizer, criterion, ):
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf # set initial "min" to infinity
    len_train = len(train_loader.dataset)
    len_valid = len(valid_loader.dataset)
    best_acc = -np.Inf
    best_sd = None

    for epoch in range(n_epochs):
        # monitor training loss
        train_loss = 0.0
        train_acc = 0.0
        valid_loss = 0.0
        valid_acc = 0.0

        #####
        # train the model #
        #####
        model.train() # prep model for training

        for data, target in train_loader:
            # clear the gradients of all optimized variables
            optimizer.zero_grad()
            # forward pass: compute predicted outputs by passing inputs to the
            output = model(data)
```

```

# calculate the loss
loss = criterion(output, target)
# backward pass: compute gradient of the loss with respect to model parameters
loss.backward()
# perform a single optimization step (parameter update)
optimizer.step()
# update running training loss
train_loss += loss.item() # *data.size(0)
# update training accuracy
num_match, _ = get_acc(output, target)
train_acc += num_match

#####
# validate the model #
#####
model.eval() # prep model for evaluation
with torch.no_grad():
    for data, target in valid_loader:
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update running validation loss
        valid_loss += loss.item()
        # update training accuracy
        num_match, _ = get_acc(output, target)
        valid_acc += num_match

train_acc = train_acc / len_train
valid_acc = valid_acc / len_valid

train_loss = train_loss / len_train * 10
valid_loss = valid_loss / len_valid * 10

# save params of the best model
if valid_acc > best_acc:
    best_acc = valid_acc
    best_sd = copy.deepcopy(model.state_dict())

if epoch % print_every == 0:
    print('Epoch: {} \tTrain Loss: {:.4f}\tValid Loss: {:.4f}\tTrain Acc: {:.4f}\tValid Acc: {:.4f}'
          .format(epoch, train_loss, valid_loss, train_acc, valid_acc))

return best_sd

```

## 4(a). Average Word2Vec vectors

In [38]:

```

# initialize the NN
model = FNN(300)
print(model)

# specify loss function (categorical cross-entropy)
criterion = nn.CrossEntropyLoss()

# specify optimizer (stochastic gradient descent) and learning rate = 0.05
optimizer = torch.optim.SGD(model.parameters(), lr=0.05)

FNN(
  (fc1): Linear(in_features=300, out_features=100, bias=True)
  (fc2): Linear(in_features=100, out_features=10, bias=True)
  (fc3): Linear(in_features=10, out_features=3, bias=True)

```

```
(dropout): Dropout(p=0.2, inplace=False)
)
```

In [39]:

```
best_sd = train(30, model, train_loader, test_loader, optimizer, criterion, p
```

```
Epoch: 1      Train Loss: 0.9719      Valid Loss: 0.8644      Train Acc: 0.4
985      Valid Acc: 0.5916
Epoch: 2      Train Loss: 0.8584      Valid Loss: 0.8527      Train Acc: 0.6
022      Valid Acc: 0.6048
Epoch: 3      Train Loss: 0.8387      Valid Loss: 0.8933      Train Acc: 0.6
128      Valid Acc: 0.5831
Epoch: 4      Train Loss: 0.8272      Valid Loss: 0.8182      Train Acc: 0.6
199      Valid Acc: 0.6283
Epoch: 5      Train Loss: 0.8194      Valid Loss: 0.8260      Train Acc: 0.6
244      Valid Acc: 0.6246
Epoch: 6      Train Loss: 0.8133      Valid Loss: 0.8310      Train Acc: 0.6
272      Valid Acc: 0.6168
Epoch: 7      Train Loss: 0.8065      Valid Loss: 0.8122      Train Acc: 0.6
315      Valid Acc: 0.6279
Epoch: 8      Train Loss: 0.8002      Valid Loss: 0.8073      Train Acc: 0.6
338      Valid Acc: 0.6306
Epoch: 9      Train Loss: 0.7949      Valid Loss: 0.8278      Train Acc: 0.6
380      Valid Acc: 0.6217
Epoch: 10     Train Loss: 0.7901      Valid Loss: 0.8030      Train Acc: 0.6
395      Valid Acc: 0.6325
Epoch: 11     Train Loss: 0.7849      Valid Loss: 0.8123      Train Acc: 0.6
435      Valid Acc: 0.6250
Epoch: 12     Train Loss: 0.7806      Valid Loss: 0.8053      Train Acc: 0.6
441      Valid Acc: 0.6327
Epoch: 13     Train Loss: 0.7759      Valid Loss: 0.8014      Train Acc: 0.6
488      Valid Acc: 0.6347
Epoch: 14     Train Loss: 0.7711      Valid Loss: 0.8129      Train Acc: 0.6
501      Valid Acc: 0.6293
Epoch: 15     Train Loss: 0.7659      Valid Loss: 0.8011      Train Acc: 0.6
509      Valid Acc: 0.6406
Epoch: 16     Train Loss: 0.7606      Valid Loss: 0.8076      Train Acc: 0.6
534      Valid Acc: 0.6359
Epoch: 17     Train Loss: 0.7572      Valid Loss: 0.8144      Train Acc: 0.6
596      Valid Acc: 0.6327
Epoch: 18     Train Loss: 0.7530      Valid Loss: 0.8047      Train Acc: 0.6
619      Valid Acc: 0.6387
Epoch: 19     Train Loss: 0.7479      Valid Loss: 0.8076      Train Acc: 0.6
616      Valid Acc: 0.6357
Epoch: 20     Train Loss: 0.7450      Valid Loss: 0.8027      Train Acc: 0.6
633      Valid Acc: 0.6388
Epoch: 21     Train Loss: 0.7381      Valid Loss: 0.8125      Train Acc: 0.6
713      Valid Acc: 0.6273
Epoch: 22     Train Loss: 0.7350      Valid Loss: 0.8501      Train Acc: 0.6
700      Valid Acc: 0.6061
Epoch: 23     Train Loss: 0.7309      Valid Loss: 0.8092      Train Acc: 0.6
730      Valid Acc: 0.6324
Epoch: 24     Train Loss: 0.7250      Valid Loss: 0.8080      Train Acc: 0.6
766      Valid Acc: 0.6336
Epoch: 25     Train Loss: 0.7213      Valid Loss: 0.8086      Train Acc: 0.6
770      Valid Acc: 0.6361
Epoch: 26     Train Loss: 0.7146      Valid Loss: 0.8282      Train Acc: 0.6
834      Valid Acc: 0.6270
Epoch: 27     Train Loss: 0.7116      Valid Loss: 0.8224      Train Acc: 0.6
798      Valid Acc: 0.6286
Epoch: 28     Train Loss: 0.7079      Valid Loss: 0.8150      Train Acc: 0.6
843      Valid Acc: 0.6324
Epoch: 29     Train Loss: 0.7034      Valid Loss: 0.8204      Train Acc: 0.6
885      Valid Acc: 0.6360
Epoch: 30     Train Loss: 0.6973      Valid Loss: 0.9531      Train Acc: 0.6
915      Valid Acc: 0.5795
```

In [40]:

```
model_best_fnn_avg = FNN(300)
model_best_fnn_avg.load_state_dict(best_sd)
```

```
_, train_acc = get_acc(model_best_fnn_avg(X_train_tensor), y_train_tensor)
_, test_acc = get_acc(model_best_fnn_avg(X_test_tensor), y_test_tensor)

print('FNN using average Word2Vec vectors:')
print('Training Accuracy =', train_acc)
print('Testing Accuracy =', test_acc)
```

FNN using average Word2Vec vectors:  
 Training Accuracy = 0.6687916666666667  
 Testing Accuracy = 0.6405833333333333

## 4(b). Concatenated Word2Vec vectors

```
In [41]: # Converting data to tensors
X_train_tensor_concat = torch.from_numpy(X_train_w2v_concat.to_numpy().astype(np.float32))
X_test_tensor_concat = torch.from_numpy(X_test_w2v_concat.to_numpy().astype(np.float32))

# Passing to DataLoader
train_tensor = data_utils.TensorDataset(X_train_tensor_concat, y_train_tensor)
train_loader = data_utils.DataLoader(train_tensor, batch_size=10, shuffle=True)

test_tensor = data_utils.TensorDataset(X_test_tensor_concat, y_test_tensor)
test_loader = data_utils.DataLoader(test_tensor, batch_size=10, shuffle=True)
```

```
In [42]: # initialize the NN
model = FNN(3000)
print(model)

# specify loss function (categorical cross-entropy)
criterion = nn.CrossEntropyLoss()

# specify optimizer (stochastic gradient descent) and learning rate = 0.01
optimizer = torch.optim.SGD(model.parameters(), lr=0.005)
```

```
FNN(
  (fc1): Linear(in_features=3000, out_features=100, bias=True)
  (fc2): Linear(in_features=100, out_features=10, bias=True)
  (fc3): Linear(in_features=10, out_features=3, bias=True)
  (dropout): Dropout(p=0.2, inplace=False)
)
```

```
In [43]: best_sd = train(20, model, train_loader, test_loader, optimizer, criterion, p
```

Epoch: 1	Train Loss: 1.0627	Valid Loss: 0.9910	Train Acc: 0.4
349	Valid Acc: 0.4943		
Epoch: 2	Train Loss: 0.9529	Valid Loss: 0.9380	Train Acc: 0.5
219	Valid Acc: 0.5373		
Epoch: 3	Train Loss: 0.9135	Valid Loss: 0.9217	Train Acc: 0.5
595	Valid Acc: 0.5558		
Epoch: 4	Train Loss: 0.8848	Valid Loss: 0.9114	Train Acc: 0.5
817	Valid Acc: 0.5657		
Epoch: 5	Train Loss: 0.8646	Valid Loss: 0.9025	Train Acc: 0.5
958	Valid Acc: 0.5694		
Epoch: 6	Train Loss: 0.8459	Valid Loss: 0.9021	Train Acc: 0.6
082	Valid Acc: 0.5660		
Epoch: 7	Train Loss: 0.8275	Valid Loss: 0.9084	Train Acc: 0.6
183	Valid Acc: 0.5675		
Epoch: 8	Train Loss: 0.8050	Valid Loss: 0.9111	Train Acc: 0.6
318	Valid Acc: 0.5700		
Epoch: 9	Train Loss: 0.7787	Valid Loss: 0.9102	Train Acc: 0.6
486	Valid Acc: 0.5653		
Epoch: 10	Train Loss: 0.7435	Valid Loss: 0.9413	Train Acc: 0.6

705	Valid Acc: 0.5614		
Epoch: 11	Train Loss: 0.7002	Valid Loss: 0.9490	Train Acc: 0.6
959	Valid Acc: 0.5589		
Epoch: 12	Train Loss: 0.6453	Valid Loss: 0.9825	Train Acc: 0.7
289	Valid Acc: 0.5584		
Epoch: 13	Train Loss: 0.5747	Valid Loss: 1.0422	Train Acc: 0.7
694	Valid Acc: 0.5488		
Epoch: 14	Train Loss: 0.4997	Valid Loss: 1.1095	Train Acc: 0.8
083	Valid Acc: 0.5472		
Epoch: 15	Train Loss: 0.4153	Valid Loss: 1.2141	Train Acc: 0.8
465	Valid Acc: 0.5347		
Epoch: 16	Train Loss: 0.3360	Valid Loss: 1.3394	Train Acc: 0.8
821	Valid Acc: 0.5331		
Epoch: 17	Train Loss: 0.2590	Valid Loss: 1.4817	Train Acc: 0.9
157	Valid Acc: 0.5302		
Epoch: 18	Train Loss: 0.1949	Valid Loss: 1.6066	Train Acc: 0.9
412	Valid Acc: 0.5248		
Epoch: 19	Train Loss: 0.1425	Valid Loss: 1.7546	Train Acc: 0.9
603	Valid Acc: 0.5288		
Epoch: 20	Train Loss: 0.1011	Valid Loss: 1.8605	Train Acc: 0.9
734	Valid Acc: 0.5292		

In [44]:

```

model_best_fnn_concat = FNN(3000)
model_best_fnn_concat.load_state_dict(best_sd)

_, train_acc = get_acc(model_best_fnn_concat(X_train_tensor_concat), y_train_)
_, test_acc = get_acc(model_best_fnn_concat(X_test_tensor_concat), y_test_ten

print('FNN using concat Word2Vec vectors:')
print('Training Accuracy =', train_acc)
print('Testing Accuracy =', test_acc)

```

FNN using concat Word2Vec vectors:  
 Training Accuracy = 0.664125  
 Testing Accuracy = 0.57

What do you conclude by comparing accuracy values you obtain with those obtained in the “Simple Models” section?

- Considering only models using Word2Vec vectors as input features, FNN (with average W2V) gives higher accuracy which is as expected since the model is more complex than using just a single perceptron.
- When considering different types of input W2V vectors for FNN, model using average W2V is better than using concatenated W2V

## 5. Recurrent Neural Networks

### 5(a). Simple RNN

```
In [45]: # Converting data to tensors
X_train_tensor_3d = torch.from_numpy(X_train_3d)
X_test_tensor_3d = torch.from_numpy(X_test_3d)

# Passing to DataLoader
train_tensor = data_utils.TensorDataset(X_train_tensor_3d, y_train_tensor)
train_loader = data_utils.DataLoader(train_tensor, batch_size=10, shuffle=False)

test_tensor = data_utils.TensorDataset(X_test_tensor_3d, y_test_tensor)
test_loader = data_utils.DataLoader(test_tensor, batch_size=10, shuffle=False)
```

```
In [46]: class model_RNN(nn.Module):
    def __init__(self, hidden_dim, n_layers):
        super(model_RNN, self).__init__()

        # Defining some parameters
        self.hidden_dim = hidden_dim
        self.n_layers = n_layers

        #Defining the layers
        # RNN Layer
        self.rnn = nn.RNN(input_size=300, hidden_size=hidden_dim, num_layers=
        # Fully connected layer
        self.fc = nn.Linear(hidden_dim, 3)

    def forward(self, x):

        batch_size = x.size(0)

        # Initializing hidden state for first input using method defined below
        hidden = self.init_hidden(batch_size)

        # Passing in the input and hidden state into the model and obtaining
        out, hidden = self.rnn(x, hidden)

        # Reshaping the outputs such that it can be fit into the fully connec
        out = out[:, -1, :]
        out = self.fc(hidden[-1])

        return out

    def init_hidden(self, batch_size):
        # This method generates the first hidden state of zeros which we'll u
        # We'll send the tensor holding the hidden state to the device we spe
        hidden = torch.zeros(self.n_layers, batch_size, self.hidden_dim)
        return hidden
```

```
In [59]: # initialize the NN
model = model_RNN(20, 2)
print(model)

# specify loss function (categorical cross-entropy)
criterion = nn.CrossEntropyLoss()

# specify optimizer (stochastic gradient descent) and learning rate = 0.01
optimizer = torch.optim.SGD(model.parameters(), lr=0.005)

model_RNN(
    (rnn): RNN(300, 20, num_layers=2, batch_first=True)
```

```
(fc): Linear(in_features=20, out_features=3, bias=True)
)
```

```
In [60]: best_sd = train(30, model, train_loader, test_loader, optimizer, criterion, p
```

Epoch: 1	Train Loss: 1.1003	Valid Loss: 1.0990	Train Acc: 0.3
327	Valid Acc: 0.3372		
Epoch: 2	Train Loss: 1.0980	Valid Loss: 1.0975	Train Acc: 0.3
436	Valid Acc: 0.3417		
Epoch: 3	Train Loss: 1.0948	Valid Loss: 1.0812	Train Acc: 0.3
575	Valid Acc: 0.4044		
Epoch: 4	Train Loss: 1.0049	Valid Loss: 0.9439	Train Acc: 0.4
827	Valid Acc: 0.5282		
Epoch: 5	Train Loss: 0.9472	Valid Loss: 0.9196	Train Acc: 0.5
256	Valid Acc: 0.5492		
Epoch: 6	Train Loss: 0.9339	Valid Loss: 0.9110	Train Acc: 0.5
381	Valid Acc: 0.5558		
Epoch: 7	Train Loss: 0.9225	Valid Loss: 0.9046	Train Acc: 0.5
465	Valid Acc: 0.5603		
Epoch: 8	Train Loss: 0.9138	Valid Loss: 0.9013	Train Acc: 0.5
513	Valid Acc: 0.5674		
Epoch: 9	Train Loss: 0.9075	Valid Loss: 0.8965	Train Acc: 0.5
560	Valid Acc: 0.5664		
Epoch: 10	Train Loss: 0.9028	Valid Loss: 0.8919	Train Acc: 0.5
593	Valid Acc: 0.5713		
Epoch: 11	Train Loss: 0.8989	Valid Loss: 0.8927	Train Acc: 0.5
614	Valid Acc: 0.5683		
Epoch: 12	Train Loss: 0.8970	Valid Loss: 0.8954	Train Acc: 0.5
649	Valid Acc: 0.5718		
Epoch: 13	Train Loss: 0.8934	Valid Loss: 0.8920	Train Acc: 0.5
679	Valid Acc: 0.5713		
Epoch: 14	Train Loss: 0.8925	Valid Loss: 0.8890	Train Acc: 0.5
678	Valid Acc: 0.5716		
Epoch: 15	Train Loss: 0.8894	Valid Loss: 0.8864	Train Acc: 0.5
699	Valid Acc: 0.5722		
Epoch: 16	Train Loss: 0.8850	Valid Loss: 0.8828	Train Acc: 0.5
737	Valid Acc: 0.5756		
Epoch: 17	Train Loss: 0.8838	Valid Loss: 0.8895	Train Acc: 0.5
749	Valid Acc: 0.5721		
Epoch: 18	Train Loss: 0.8814	Valid Loss: 0.8822	Train Acc: 0.5
767	Valid Acc: 0.5833		
Epoch: 19	Train Loss: 0.8794	Valid Loss: 0.8821	Train Acc: 0.5
785	Valid Acc: 0.5797		
Epoch: 20	Train Loss: 0.8783	Valid Loss: 0.8774	Train Acc: 0.5
799	Valid Acc: 0.5777		
Epoch: 21	Train Loss: 0.8764	Valid Loss: 0.8862	Train Acc: 0.5
820	Valid Acc: 0.5832		
Epoch: 22	Train Loss: 0.8741	Valid Loss: 0.8802	Train Acc: 0.5
834	Valid Acc: 0.5829		
Epoch: 23	Train Loss: 0.8734	Valid Loss: 0.8797	Train Acc: 0.5
859	Valid Acc: 0.5746		
Epoch: 24	Train Loss: 0.8728	Valid Loss: 0.8825	Train Acc: 0.5
861	Valid Acc: 0.5778		
Epoch: 25	Train Loss: 0.9082	Valid Loss: 0.9653	Train Acc: 0.5
593	Valid Acc: 0.5389		
Epoch: 26	Train Loss: 0.9248	Valid Loss: 0.9105	Train Acc: 0.5
509	Valid Acc: 0.5647		
Epoch: 27	Train Loss: 0.9318	Valid Loss: 0.9223	Train Acc: 0.5
460	Valid Acc: 0.5498		
Epoch: 28	Train Loss: 0.9029	Valid Loss: 0.9365	Train Acc: 0.5
729	Valid Acc: 0.5465		
Epoch: 29	Train Loss: 0.9028	Valid Loss: 0.8848	Train Acc: 0.5
737	Valid Acc: 0.5885		
Epoch: 30	Train Loss: 0.8874	Valid Loss: 0.8896	Train Acc: 0.5
884	Valid Acc: 0.5883		

```
In [61]: model_best_rnn = model_RNN(20, 2)
model_best_rnn.load_state_dict(best_sd)
```



```
_, train_acc = get_acc(model_best_rnn(X_train_tensor_3d), y_train_tensor)
_, test_acc = get_acc(model_best_rnn(X_test_tensor_3d), y_test_tensor)

print('RNN:')
print('Training Accuracy =', train_acc)
print('Testing Accuracy =', test_acc)
```

RNN:  
 Training Accuracy = 0.5914583333333333  
 Testing Accuracy = 0.5885

## 5(b). GRU: Gated Recurrent Unit

In [50]:

```
class model_GRU(nn.Module):
    def __init__(self, hidden_dim, n_layers):
        super(model_GRU, self).__init__()

        # Defining some parameters
        self.hidden_dim = hidden_dim
        self.n_layers = n_layers

        #Defining the layers
        # RNN Layer
        self.rnn = nn.GRU(input_size=300, hidden_size=hidden_dim, num_layers=
        # Fully connected layer
        self.fc = nn.Linear(hidden_dim, 3)

    def forward(self, x):

        batch_size = x.size(0)

        # Initializing hidden state for first input using method defined below
        hidden = self.init_hidden(batch_size)

        # Passing in the input and hidden state into the model and obtaining
        out, _ = self.rnn(x, hidden)

        # Reshaping the outputs such that it can be fit into the fully connec
        out = out[:, -1, :]
        out = self.fc(out)

        return out

    def init_hidden(self, batch_size):
        # This method generates the first hidden state of zeros which we'll u
        # We'll send the tensor holding the hidden state to the device we spe
        hidden = torch.zeros(self.n_layers, batch_size, self.hidden_dim)
        return hidden
```

In [51]:

```
# initialize the NN
model = model_GRU(20, 2)
print(model)

# specify loss function (categorical cross-entropy)
criterion = nn.CrossEntropyLoss()

# specify optimizer (stochastic gradient descent) and learning rate = 0.01
optimizer = torch.optim.SGD(model.parameters(), lr=0.05)

model_GRU(
    (rnn): GRU(300, 20, num_layers=2, batch_first=True)
```



```
(fc): Linear(in_features=20, out_features=3, bias=True)
)
```

In [52]:

```
best_sd = train(30, model, train_loader, test_loader, optimizer, criterion, p
```

```
Epoch: 1      Train Loss: 1.0956      Valid Loss: 1.0796      Train Acc: 0.3
530      Valid Acc: 0.3924
Epoch: 2      Train Loss: 0.9740      Valid Loss: 0.8967      Train Acc: 0.5
028      Valid Acc: 0.5582
Epoch: 3      Train Loss: 0.9000      Valid Loss: 0.8669      Train Acc: 0.5
615      Valid Acc: 0.5842
Epoch: 4      Train Loss: 0.8690      Valid Loss: 0.8438      Train Acc: 0.5
881      Valid Acc: 0.5987
Epoch: 5      Train Loss: 0.8463      Valid Loss: 0.8308      Train Acc: 0.6
048      Valid Acc: 0.6096
Epoch: 6      Train Loss: 0.8312      Valid Loss: 0.8206      Train Acc: 0.6
156      Valid Acc: 0.6178
Epoch: 7      Train Loss: 0.8196      Valid Loss: 0.8127      Train Acc: 0.6
219      Valid Acc: 0.6248
Epoch: 8      Train Loss: 0.8102      Valid Loss: 0.8071      Train Acc: 0.6
273      Valid Acc: 0.6273
Epoch: 9      Train Loss: 0.8025      Valid Loss: 0.8028      Train Acc: 0.6
317      Valid Acc: 0.6278
Epoch: 10     Train Loss: 0.7959      Valid Loss: 0.7995      Train Acc: 0.6
341      Valid Acc: 0.6298
Epoch: 11     Train Loss: 0.7902      Valid Loss: 0.7967      Train Acc: 0.6
374      Valid Acc: 0.6312
Epoch: 12     Train Loss: 0.7849      Valid Loss: 0.7944      Train Acc: 0.6
405      Valid Acc: 0.6330
Epoch: 13     Train Loss: 0.7801      Valid Loss: 0.7924      Train Acc: 0.6
433      Valid Acc: 0.6346
Epoch: 14     Train Loss: 0.7756      Valid Loss: 0.7907      Train Acc: 0.6
467      Valid Acc: 0.6358
Epoch: 15     Train Loss: 0.7713      Valid Loss: 0.7892      Train Acc: 0.6
489      Valid Acc: 0.6368
Epoch: 16     Train Loss: 0.7672      Valid Loss: 0.7880      Train Acc: 0.6
519      Valid Acc: 0.6372
Epoch: 17     Train Loss: 0.7633      Valid Loss: 0.7869      Train Acc: 0.6
542      Valid Acc: 0.6377
Epoch: 18     Train Loss: 0.7595      Valid Loss: 0.7861      Train Acc: 0.6
562      Valid Acc: 0.6379
Epoch: 19     Train Loss: 0.7558      Valid Loss: 0.7855      Train Acc: 0.6
581      Valid Acc: 0.6385
Epoch: 20     Train Loss: 0.7522      Valid Loss: 0.7850      Train Acc: 0.6
599      Valid Acc: 0.6395
Epoch: 21     Train Loss: 0.7487      Valid Loss: 0.7848      Train Acc: 0.6
622      Valid Acc: 0.6395
Epoch: 22     Train Loss: 0.7452      Valid Loss: 0.7847      Train Acc: 0.6
641      Valid Acc: 0.6396
Epoch: 23     Train Loss: 0.7418      Valid Loss: 0.7848      Train Acc: 0.6
662      Valid Acc: 0.6394
Epoch: 24     Train Loss: 0.7385      Valid Loss: 0.7850      Train Acc: 0.6
676      Valid Acc: 0.6401
Epoch: 25     Train Loss: 0.7352      Valid Loss: 0.7853      Train Acc: 0.6
696      Valid Acc: 0.6397
Epoch: 26     Train Loss: 0.7319      Valid Loss: 0.7858      Train Acc: 0.6
703      Valid Acc: 0.6395
Epoch: 27     Train Loss: 0.7287      Valid Loss: 0.7863      Train Acc: 0.6
720      Valid Acc: 0.6387
Epoch: 28     Train Loss: 0.7255      Valid Loss: 0.7869      Train Acc: 0.6
741      Valid Acc: 0.6390
Epoch: 29     Train Loss: 0.7223      Valid Loss: 0.7876      Train Acc: 0.6
755      Valid Acc: 0.6372
Epoch: 30     Train Loss: 0.7192      Valid Loss: 0.7884      Train Acc: 0.6
780      Valid Acc: 0.6362
```

In [53]:

```
model_best_gru = model_GRU(20, 2)
model_best_gru.load_state_dict(best_sd)
```

```
_, train_acc = get_acc(model_best_gru(X_train_tensor_3d), y_train_tensor)
_, test_acc = get_acc(model_best_gru(X_test_tensor_3d), y_test_tensor)

print('GRU:')
print('Training Accuracy =', train_acc)
print('Testing Accuracy =', test_acc)
```

GRU:  
 Training Accuracy = 0.6696041666666667  
 Testing Accuracy = 0.6400833333333333

## 5(c). LSTM

In [54]:

```
class model_LSTM(nn.Module):
    def __init__(self, hidden_dim, n_layers):
        super(model_LSTM, self).__init__()

        # Defining some parameters
        self.hidden_dim = hidden_dim
        self.n_layers = n_layers

        #Defining the layers
        # RNN Layer
        self.rnn = nn.LSTM(input_size=300, hidden_size=hidden_dim, num_layers=
        # Fully connected layer
        self.fc = nn.Linear(hidden_dim, 3)

    def forward(self, x):

        batch_size = x.size(0)

        # Initializing hidden state for first input using method defined below
        h0, c0 = self.init_hidden(batch_size)

        # Passing in the input and hidden state into the model and obtaining
        out, _ = self.rnn(x, (h0, c0))

        # Reshaping the outputs such that it can be fit into the fully connect
        out = out[:, -1, :]
        out = self.fc(out)

        return out

    def init_hidden(self, batch_size):
        # This method generates the first hidden state of zeros which we'll use
        # We'll send the tensor holding the hidden state to the device we spe
        h0 = torch.randn(self.n_layers, batch_size, self.hidden_dim)
        c0 = torch.randn(self.n_layers, batch_size, self.hidden_dim)
        return h0, c0
```

In [55]:

```
# initialize the NN
model = model_LSTM(20, 2)
print(model)

# specify loss function (categorical cross-entropy)
criterion = nn.CrossEntropyLoss()

# specify optimizer (stochastic gradient descent) and learning rate = 0.01
optimizer = torch.optim.SGD(model.parameters(), lr=0.03)
```

```

model_LSTM(
    (rnn): LSTM(300, 20, num_layers=2, batch_first=True)
    (fc): Linear(in_features=20, out_features=3, bias=True)
)

```

```
In [56]: best_sd = train(30, model, train_loader, test_loader, optimizer, criterion, p
```

Epoch: 1	Train Loss: 1.0992	Valid Loss: 1.0988	Train Acc: 0.3
311	Valid Acc: 0.3333		
Epoch: 2	Train Loss: 1.0981	Valid Loss: 1.0959	Train Acc: 0.3
412	Valid Acc: 0.3375		
Epoch: 3	Train Loss: 1.0347	Valid Loss: 0.9716	Train Acc: 0.4
532	Valid Acc: 0.5158		
Epoch: 4	Train Loss: 0.9684	Valid Loss: 0.9413	Train Acc: 0.5
137	Valid Acc: 0.5275		
Epoch: 5	Train Loss: 0.9424	Valid Loss: 0.9254	Train Acc: 0.5
326	Valid Acc: 0.5370		
Epoch: 6	Train Loss: 0.9230	Valid Loss: 0.8986	Train Acc: 0.5
487	Valid Acc: 0.5666		
Epoch: 7	Train Loss: 0.9079	Valid Loss: 0.8851	Train Acc: 0.5
671	Valid Acc: 0.5785		
Epoch: 8	Train Loss: 0.8925	Valid Loss: 0.8725	Train Acc: 0.5
758	Valid Acc: 0.5877		
Epoch: 9	Train Loss: 0.8816	Valid Loss: 0.8616	Train Acc: 0.5
830	Valid Acc: 0.5968		
Epoch: 10	Train Loss: 0.8729	Valid Loss: 0.8581	Train Acc: 0.5
909	Valid Acc: 0.5996		
Epoch: 11	Train Loss: 0.8665	Valid Loss: 0.8529	Train Acc: 0.5
940	Valid Acc: 0.6028		
Epoch: 12	Train Loss: 0.8591	Valid Loss: 0.8476	Train Acc: 0.5
978	Valid Acc: 0.6023		
Epoch: 13	Train Loss: 0.8524	Valid Loss: 0.8433	Train Acc: 0.6
020	Valid Acc: 0.6062		
Epoch: 14	Train Loss: 0.8480	Valid Loss: 0.8415	Train Acc: 0.6
031	Valid Acc: 0.6112		
Epoch: 15	Train Loss: 0.8424	Valid Loss: 0.8362	Train Acc: 0.6
072	Valid Acc: 0.6129		
Epoch: 16	Train Loss: 0.8376	Valid Loss: 0.8377	Train Acc: 0.6
108	Valid Acc: 0.6113		
Epoch: 17	Train Loss: 0.8332	Valid Loss: 0.8318	Train Acc: 0.6
146	Valid Acc: 0.6125		
Epoch: 18	Train Loss: 0.8307	Valid Loss: 0.8308	Train Acc: 0.6
138	Valid Acc: 0.6151		
Epoch: 19	Train Loss: 0.8261	Valid Loss: 0.8311	Train Acc: 0.6
173	Valid Acc: 0.6107		
Epoch: 20	Train Loss: 0.8232	Valid Loss: 0.8274	Train Acc: 0.6
183	Valid Acc: 0.6150		
Epoch: 21	Train Loss: 0.8191	Valid Loss: 0.8293	Train Acc: 0.6
205	Valid Acc: 0.6119		
Epoch: 22	Train Loss: 0.8165	Valid Loss: 0.8233	Train Acc: 0.6
230	Valid Acc: 0.6143		
Epoch: 23	Train Loss: 0.8115	Valid Loss: 0.8193	Train Acc: 0.6
262	Valid Acc: 0.6166		
Epoch: 24	Train Loss: 0.8096	Valid Loss: 0.8200	Train Acc: 0.6
284	Valid Acc: 0.6213		
Epoch: 25	Train Loss: 0.8069	Valid Loss: 0.8190	Train Acc: 0.6
291	Valid Acc: 0.6191		
Epoch: 26	Train Loss: 0.8051	Valid Loss: 0.8164	Train Acc: 0.6
305	Valid Acc: 0.6252		
Epoch: 27	Train Loss: 0.7997	Valid Loss: 0.8219	Train Acc: 0.6
317	Valid Acc: 0.6158		
Epoch: 28	Train Loss: 0.7972	Valid Loss: 0.8184	Train Acc: 0.6
322	Valid Acc: 0.6203		
Epoch: 29	Train Loss: 0.7956	Valid Loss: 0.8148	Train Acc: 0.6
350	Valid Acc: 0.6215		
Epoch: 30	Train Loss: 0.7926	Valid Loss: 0.8166	Train Acc: 0.6
354	Valid Acc: 0.6200		

In [58]:

```
model_best_lstm = model_LSTM(20, 2)
model_best_lstm.load_state_dict(best_sd)

_, train_acc = get_acc(model_best_lstm(X_train_tensor_3d), y_train_tensor)
_, test_acc = get_acc(model_best_lstm(X_test_tensor_3d), y_test_tensor)

print('LSTM:')
print('Training Accuracy =', train_acc)
print('Testing Accuracy =', test_acc)
```

LSTM:

Training Accuracy = 0.6322708333333333

Testing Accuracy = 0.6245

What do you conclude by comparing accuracy values you obtain by GRU, LSTM, and simple RNN?

- GRU gives the best test accuracy at 64% while LSTM's accuracy is slightly lower at 62.5%. For simple RNN, the best test accuracy achieved is at 58.9%. This can be inferred that this sentiment prediction task also relies on the long-term dependency of the input reviews, and LSTM and GRU are doing better task in storing and retrieving long-term dependency than simple RNN

## References

- [https://radimrehurek.com/gensim/auto\\_examples/tutorials/run\\_word2vec.html](https://radimrehurek.com/gensim/auto_examples/tutorials/run_word2vec.html)
- <https://www.kaggle.com/code/mishra1993/pytorch-multi-layer-perceptron-mnist/notebook>
- <https://stackoverflow.com/questions/3438756/some-built-in-to-pad-a-list-in-python>
- <https://blog.floydhub.com/a-beginners-guide-on-recurrent-neural-networks-with-pytorch/>
- [https://androidkt.com/copy-pytorch-model-using-deepcopy-and-state\\_dict/](https://androidkt.com/copy-pytorch-model-using-deepcopy-and-state_dict/)